

compboost: Modular Framework for Component-Wise Boosting

04 July 2018

Summary

Component-wise boosting applies the boosting framework to statistical models, e.g., general additive models using component-wise smoothing splines (Schmid and Hothorn 2008). Boosting these kinds of models maintains interpretability and enables unbiased model selection in high dimensional feature spaces.

The R (Team 2016) package `compboost` is an implementation of component-wise boosting written in C++ using `Armadillo` (Sanderson and Curtin 2016) to obtain high runtime performance and full memory control. The main idea is to provide a modular class system which can be extended without editing the source code. Therefore, it is possible to use R functions as well as C++ functions for custom base-learners, losses, logging mechanisms or stopping criteria.

In addition to tree based boosting implementations as `xgboost` (Chen et al. 2018), `compboost`, which is not a tree based method, maintains interpretability by estimating parameter for each used base-learner. This allows visualizing the selected effects, jumping back and forth in the algorithm, and looking into the model how the learners are selected to obtain information about the feature importance.

How to Use

The package provides two high level wrapper functions `boostLinear()` and `boostSplines()` to boost linear models or general additive models using p-splines of each numerical feature. The data used for the demo are from the `mlbench` (Leisch and Dimitriadou 2010) package:

```
library(compboost)

# Load data set with binary classification task:
data(PimaIndiansDiabetes, package = "mlbench")

# Quadratic loss as ordinary regression loss:
cboost = boostSplines(data = PimaIndiansDiabetes, target = "diabetes",
  loss = BinomialLoss$new())
```

The resulting model is an R6 (Chang 2017) object. Hence, `mod` has member functions to access the elements of the model such as the names of registered base-learner, selected base-learner, the estimated parameter, or to continue the training:

```
cboost$getBaselearnerNames()
## [1] "pregnant_spline" "glucose_spline"  "pressure_spline" "triceps_spline"
## [5] "insulin_spline"  "mass_spline"     "pedigree_spline" "age_spline"

selected.features = mod$selected()
table(selected.features)
## selected.features
##    age_spline glucose_spline    mass_spline
##          23             61             16
```




Figure 1: Visualize compboost

```
params = cboost$coef()
str(params)
## List of 4
## $ age_spline : num [1:24, 1] 0.40127 0.25655 0.14807 0.11766 -0.00586 ...
## $ glucose_spline: num [1:24, 1] -0.2041 0.0343 0.2703 0.4921 0.6856 ...
## $ mass_spline : num [1:24, 1] 0.0681 0.0949 0.1216 0.1473 0.1714 ...
## $ offset : num 0.312

cboost$train(3000)
##
## You have already trained 100 iterations.
## Train 2900 additional iterations.
##
```

Additionally, it is possible to visualize the effect of single features by calling the member function `plot()` and passing the name of a specific learner. Furthermore, a vector of iterations can be used to plot the effect at different stages of the model:

```
cboost$plot("age_spline", iters = c(100, 500, 1000, 2000, 3000))
```

Instead of using `boostLinear()` or `boostSplines()` one can also explicitly define the parts of the algorithm by using the R6 interface:

```
cboost = Compboost$new(data = PimaIndiansDiabetes, target = "diabetes",
loss = BinomialLoss$new())

# Adding a linear and spline base-learner to the Compboost object:
cboost$addBaselearner(feature = "mass", id = "linear", PolynomialBlearner,
degree = 1, intercept = TRUE)
cboost$addBaselearner(feature = "age", id = "spline", PSplineBlearner,
degree = 3, n.knots = 10, penalty = 2, differences = 2)

cboost$train(2000, trace = FALSE)
```

```

cboost
## Component-Wise Gradient Boosting
##
## Trained on PimaIndiansDiabetes with target diabetes
## Number of base-learners: 2
## Learning rate: 0.05
## Iterations: 2000
## Positive class: neg
## Offset: 0.3118
##
## BinomialLoss Loss:
##
## Loss function:  $L(y, x) = \log(1 + \exp(-2yf(x)))$ 
##
##

```

A similar software is the well known R implementation **mboost** (Hothorn et al. 2017). The advantage of **mboost** over **compboost** is the extensive functionality which includes more base-learners and loss functions (families). Nevertheless, **mboost** has issues if trained on large datasets. In addition, **compboost** is much faster in terms of runtime and uses much less memory. This makes **compboost** more applicable for big data.

The modular principle of **compboost** allows to extend the algorithm to do more complicated analyses as boosting functional data, investigating on different optimizer, or improve the intrinsic feature selection using resampling.

References

- Chang, Winston. 2017. *R6: Classes with Reference Semantics*. <https://CRAN.R-project.org/package=R6>.
- Chen, Tianqi, Tong He, Michael Benesty, Vadim Khotilovich, Yuan Tang, Hyunsu Cho, Kailong Chen, et al. 2018. *Xgboost: Extreme Gradient Boosting*. <https://CRAN.R-project.org/package=xgboost>.
- Hothorn, Torsten, Peter Buehlmann, Thomas Kneib, Matthias Schmid, and Benjamin Hofner. 2017. *mboost: Model-Based Boosting*. <https://CRAN.R-project.org/package=mboost>.
- Leisch, Friedrich, and Evgenia Dimitriadou. 2010. *Mlbench: Machine Learning Benchmark Problems*.
- Sanderson, Conrad, and Ryan Curtin. 2016. “Armadillo: A Template-Based C++ Library for Linear Algebra.” *Journal of Open Source Software* 1 (2). Journal of Open Source Software: 26. <https://doi.org/10.21105/joss.00026>.
- Schmid, Matthias, and Torsten Hothorn. 2008. “Boosting Additive Models Using Component-Wise P-Splines.” *Computational Statistics & Data Analysis* 53 (2). Elsevier: 298–311.
- Team, R Core. 2016. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.